

# TimeSeries2021\_Mar9\_Forecasting

March 10, 2021

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
```

```
[3]: df = pd.read_excel("indicator8_2014_1.xlsx")
```

```
[4]: df.head()
```

```
[4]: Average Global Temperature, 1880-2013      Unnamed: 1  Unnamed: 2  \
0      NaN      NaN      NaN
1      Year      Temperature      NaN
2      NaN      Degrees Fahrenheit      NaN
3      NaN      NaN      NaN
4      1880      56.822      NaN

      Unnamed: 3  Unnamed: 4
0      NaN      NaN
1      NaN      NaN
2      NaN      NaN
3      NaN      NaN
4      NaN      NaN
```

```
[5]: df.tail()
```

```
[5]: Average Global Temperature, 1880-2013      Unnamed: 1  Unnamed: 2  \
136      2012      58.244      NaN
137      2013      58.298      NaN
138      NaN      NaN      NaN
139 Source: Compiled by Earth Policy Institute fro...      NaN      NaN
140      NaN      NaN      NaN

      Unnamed: 3  Unnamed: 4
136      NaN      NaN
137      NaN      NaN
138      NaN      NaN
139      NaN      NaN
```

140           NaN           NaN

```
[6]: # Add parameters to fix various problems.  
df = pd.read_excel("indicator8_2014_1.xlsx", index_col=0, usecols="A:B",  
                  skiprows=4, skipfooter=3,  
                  parse_dates=True, names=["Date", "Temp"])
```

```
[7]: df.head()
```

```
[7]:
```

	Date	Temp
	1880-01-01	56.822
	1881-01-01	56.984
	1882-01-01	56.912
	1883-01-01	56.876
	1884-01-01	56.732


```
[8]: df.tail()
```

```
[8]:
```

	Date	Temp
	2009-01-01	58.280
	2010-01-01	58.406
	2011-01-01	58.190
	2012-01-01	58.244
	2013-01-01	58.298

```
[9]: df.plot()
```

```
[9]: <matplotlib.axes._subplots.AxesSubplot at 0x11d0f8f10>
```



```
[10]: type(df.index[0])
```

```
[10]: pandas._libs.tslibs.timestamps.Timestamp
```

Rather than treat the average temperature over a particular year as happening at a particular time stamp, instead we could consider the average temperature to be happening over a particular *time period*. Python can distinguish between the two.

```
[11]: df.index = df.index.to_period('Y')
```

```
[13]: type(df.index[0])
```

```
[13]: pandas._libs.tslibs.period.Period
```


```
[14]: df.head()
```

```
[14]:
```

	Temp
Date	
1880	56.822
1881	56.984
1882	56.912
1883	56.876
1884	56.732

```
[15]: df.plot()
```

```
[15]: <matplotlib.axes._subplots.AxesSubplot at 0x111a49e10>
```



```
[16]: from statsmodels.tsa.holtwinters import ExponentialSmoothing
```

```
[17]: model = ExponentialSmoothing(df)
```

```
[18]: # This is simple exponential smoothing.
fit = model.fit()
fit.summary()
```

```
[18]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

```

                        ExponentialSmoothing Model Results
=====
Dep. Variable:                endog    No. Observations:                134
Model:                        ExponentialSmoothing    SSE                    3.828
Optimized:                    True      AIC                     -472.428
Trend:                        None      BIC                     -466.632
Seasonal:                    None      AICC                    -472.118
Seasonal Periods:            None      Date:                    Wed, 10 Mar 2021
Box-Cox:                     False     Time:                    23:15:09
Box-Cox Coeff.:              None
=====
                        coeff                code                optimized
=====
```

```
-----
smoothing_level          0.4601695          alpha          True
initial_level            56.871623          1.0             True
-----
```


```
"""
```

```
[19]: dfhat = fit.predict(0,len(df)+20)
```

```
[20]: plt.figure()
      df.plot()
      dfhat.plot()
```

```
[20]: <matplotlib.axes._subplots.AxesSubplot at 0x1261986d0>
```

```
<Figure size 432x288 with 0 Axes>
```



```
[21]: # This is Holt with trend (also called double exponential smoothing).
      # But here the best fit turned out to be a constant trend.
      model = ExponentialSmoothing(df, trend="additive")
      fit = model.fit()
      fit.summary()
```

```
[21]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

# ExponentialSmoothing Model Results

```
=====
Dep. Variable:                endog    No. Observations:      134
Model:                    ExponentialSmoothing    SSE              3.750
Optimized:                  True    AIC              -471.195
Trend:                      Additive    BIC              -459.604
Seasonal:                   None    AICC             -470.534
Seasonal Periods:          None    Date:            Wed, 10 Mar 2021
Box-Cox:                   False    Time:            23:15:15
Box-Cox Coeff.:           None
=====
```


	coeff	code	optimized
smoothing_level	0.4156884	alpha	True
smoothing_slope	0.000000	beta	True
initial_level	56.843829	l.0	True
initial_slope	0.0107676	b.0	True

"""

```
[22]: dfhat = fit.predict(0,len(df)+20)
plt.figure()
df.plot()
dfhat.plot()
```

```
[22]: <matplotlib.axes._subplots.AxesSubplot at 0x12614e650>
```

```
<Figure size 432x288 with 0 Axes>
```



```
[23]: # The "Pyramid ARIMA" package has a R-style "auto.arima" function. To install:
# pip install pmdarima
import pmdarima as pm
```

```
[24]: fit = pm.auto_arima(df) # auto_arima returns a fitted model
```

```
[25]: fit.summary()
```

```
[25]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

#### SARIMAX Results

```
=====
Dep. Variable:          y      No. Observations:          134
Model:                SARIMAX(1, 1, 3)      Log Likelihood          53.392
Date:                Wed, 10 Mar 2021      AIC              -94.784
Time:                23:15:32              BIC              -77.441
Sample:              0                  HQIC              -87.736
                  - 134
Covariance Type:      opg
=====
```

	coef	std err	z	P> z	[0.025	0.975]
intercept	0.0216	0.010	2.129	0.033	0.002	0.041
ar.L1	-0.9478	0.063	-14.969	0.000	-1.072	-0.824
ma.L1	0.5644	0.112	5.054	0.000	0.346	0.783

```

ma.L2          -0.5824      0.095      -6.154      0.000      -0.768      -0.397
ma.L3          -0.2821      0.091      -3.101      0.002      -0.460      -0.104
sigma2         0.0261      0.004       6.888      0.000      0.019      0.034
=====
===
Ljung-Box (Q):                39.46   Jarque-Bera (JB):
1.79
Prob(Q):                      0.49   Prob(JB):
0.41
Heteroskedasticity (H):       1.55   Skew:
-0.09
Prob(H) (two-sided):          0.15   Kurtosis:
2.46
=====
===

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-
step).
"""

```

```
[26]: forecasts = pd.Series(fit.predict(100), index=pd.period_range('2014',
↪periods=100, freq='Y'))
```

```
[27]: forecasts.head()
```


```
[27]: 2014      58.287975
      2015      58.319309
      2016      58.317231
      2017      58.340785
      2018      58.340045
      Freq: A-DEC, dtype: float64
```

```
[28]: plt.figure()
      df.plot()
      forecasts.plot()
```

```
[28]: <matplotlib.axes._subplots.AxesSubplot at 0x12a0d3050>
```

```
<Figure size 432x288 with 0 Axes>
```





```
[29]: # We can also use "predict" to return confidence intervals.
fit.predict(100, return_conf_int=True)
```

```
[29]: (array([58.28797495, 58.31930892, 58.31723148, 58.34078499, 58.34004505,
58.36233085, 58.36279247, 58.38393942, 58.38548047, 58.40560432,
58.40811507, 58.42731981, 58.43070171, 58.44908077, 58.45324527,
58.47088256, 58.47575011, 58.49272105, 58.49822019, 58.51459249,
58.52065903, 58.53649354, 58.5430698 , 58.55842119, 58.56545536,
58.58037273, 58.58781827, 58.60234574, 58.61016084, 58.62433803,
58.63248513, 58.64634765, 58.654793 , 58.66837284, 58.67708611,
58.690412 , 58.69936598, 58.71246373, 58.72163394, 58.73452673,
58.7438912 , 58.75659988, 58.76613886, 58.77868213, 58.78837789,
58.80077257, 58.81060916, 58.82287035, 58.83283346, 58.84497473,
58.85505151, 58.86708505, 58.87726394, 58.88920069, 58.89947131,
58.91132113, 58.92167415, 58.93344586, 58.94387291, 58.95557445,
58.96606802, 58.97770652, 58.98825983, 58.9998417 , 59.01044868,
59.02197968, 59.03263489, 59.04412018, 59.0548187 , 59.06626295,
59.07700038, 59.08840774, 59.09918013, 59.11055435, 59.12135815,
59.13270261, 59.14353462, 59.15485233, 59.16570969, 59.17700338,
59.18788351, 59.19915562, 59.2100562 , 59.22130892, 59.23222788,
59.24346319, 59.25439865, 59.26561831, 59.27656861, 59.28777421,
59.29873783, 59.30993081, 59.32090639, 59.33208803, 59.34307436,
59.35424581, 59.3652418 , 59.37640409, 59.38740876, 59.39856283]),
array([[57.9713109 , 58.604639 ]],
```

[57.94728736, 58.69133047],  
[57.92448469, 58.70997826],  
[57.93493391, 58.74663606],  
[57.91547875, 58.76461136],  
[57.92533913, 58.79932258],  
[57.90866846, 58.81691649],  
[57.91794381, 58.84993503],  
[57.90363592, 58.86732501],  
[57.91235465, 58.89885399],  
[57.90008019, 58.91614995],  
[57.90828253, 58.9463571 ],  
[57.89777642, 58.96362701],  
[57.90550756, 58.99265398],  
[57.8965516 , 59.00993894],  
[57.90385787, 59.03790726],  
[57.89626941, 59.05523082],  
[57.90319613, 59.08224597],  
[57.89682026, 59.09962012],  
[57.90341055, 59.12577443],  
[57.89811454, 59.14320351],  
[57.90440872, 59.16857836],  
[57.9000779 , 59.1860617 ],  
[57.90611321, 59.21072916],  
[57.90264785, 59.22826287],  
[57.90845841, 59.25228705],  
[57.90577128, 59.26986526],  
[57.91138811, 59.29330336],  
[57.90940262, 59.31091907],  
[57.9148538 , 59.33382227],  
[57.91350236, 59.3514679 ],  
[57.9188132 , 59.3738821 ],  
[57.91803603, 59.39154996],  
[57.92322927, 59.4135164 ],  
[57.92297331, 59.43119891],  
[57.92806929, 59.45275471],  
[57.92828736, 59.4704446 ],  
[57.93330424, 59.49162322],  
[57.93395425, 59.50931363],  
[57.93890821, 59.53014526],  
[57.93995257, 59.54782984],  
[57.94485798, 59.56834178],  
[57.94626304, 59.58601468],  
[57.95113264, 59.60623163],  
[57.95286823, 59.62388754],  
[57.95771328, 59.64383186],  
[57.95975232, 59.661466 ],  
[57.96458273, 59.68115796],

[57.96690087, 59.69876606],  
[57.97172538, 59.71822408],  
[57.97430067, 59.73580235],  
[57.97912694, 59.75504315],  
[57.98193961, 59.77258826],  
[57.98677434, 59.79162705],  
[57.98980652, 59.8091361 ],  
[57.99465553, 59.82798672],  
[57.99789109, 59.84545721],  
[58.00275945, 59.86413227],  
[58.00618378, 59.88156205],  
[58.01107586, 59.90007305],  
[58.01467571, 59.91746032],  
[58.01959528, 59.93581775],  
[58.02335865, 59.95316101],  
[58.02830893, 59.97137447],  
[58.0322249 , 59.98867247],  
[58.03720863, 60.00675074],  
[58.04126727, 60.0240025 ],  
[58.04628676, 60.04195361],  
[58.05047904, 60.05915837],  
[58.05553622, 60.07698967],  
[58.0598539 , 60.09414686],  
[58.06495035, 60.11186513],  
[58.06938591, 60.12897434],  
[58.07452292, 60.14658578],  
[58.0790695 , 60.1636468 ],  
[58.0842481 , 60.18115711],  
[58.08889941, 60.19816982],  
[58.0941204 , 60.21558426],  
[58.09887067, 60.2325487 ],  
[58.10413465, 60.24987212],  
[58.1089786 , 60.26678841],  
[58.11428598, 60.28402526],  
[58.11921876, 60.30089364],  
[58.12456979, 60.31804806],  
[58.12958693, 60.33486883],  
[58.13498174, 60.35194463],  
[58.14007914, 60.36871817],  
[58.14551774, 60.38571889],  
[58.15069159, 60.40244563],  
[58.15617388, 60.41937455],  
[58.16142069, 60.43605497],  
[58.16694647, 60.45291514],  
[58.172263 , 60.46954979],  
[58.17783202, 60.48634403],  
[58.18321527, 60.50293346],

```
[58.18882719, 60.51966442],
[58.19427438, 60.53620923],
[58.19992883, 60.55287935],
[58.20543736, 60.56938017],
[58.2111339 , 60.58599175]]))
```

```
[30]: forecasts = pd.DataFrame(fit.predict(100, return_conf_int=True)[1],
                               index=pd.period_range('2014', periods=100, freq='Y'))
```

```
[31]: forecasts.head()
```

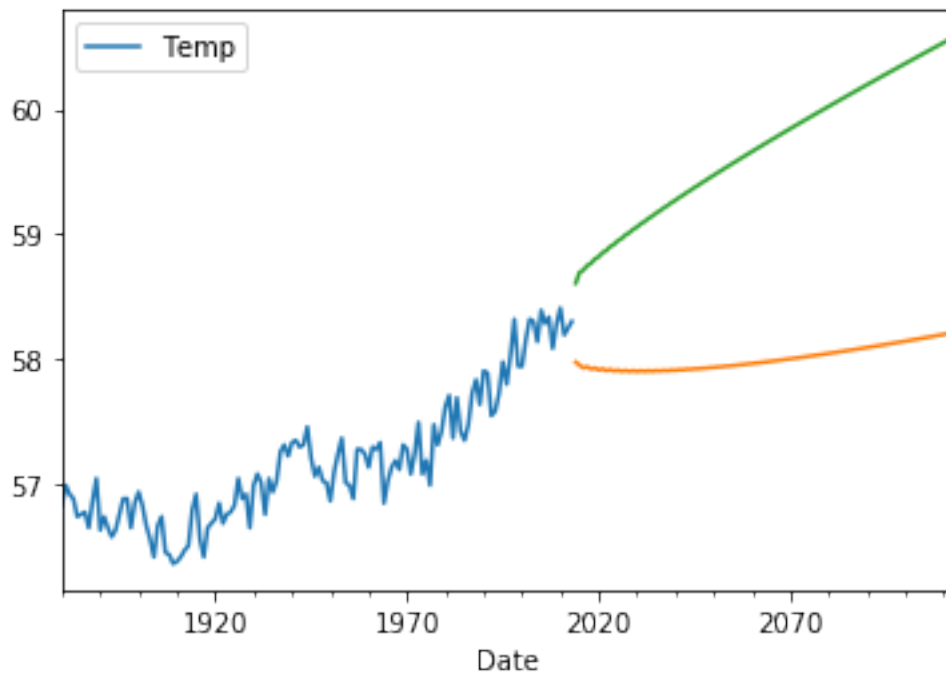
```
[31]:
```

	0	1
2014	57.971311	58.604639
2015	57.947287	58.691330
2016	57.924485	58.709978
2017	57.934934	58.746636
2018	57.915479	58.764611

```
[32]: plt.figure()
df.plot()
forecasts[0].plot()
forecasts[1].plot()
```

```
[32]: <matplotlib.axes._subplots.AxesSubplot at 0x12a0845d0>
```

<Figure size 432x288 with 0 Axes>



```
[33]: # Let's fit a trend stationary rather than difference stationary model.
fit = pm.auto_arima(df, d=0, suppress_warnings=True, trend="ct")
fit.summary()
```

```
[33]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

```

                                SARIMAX Results
=====
Dep. Variable:                  y      No. Observations:                  134
Model:                        SARIMAX(3, 0, 3)  Log Likelihood                  53.958
Date:                        Wed, 10 Mar 2021  AIC                      -89.916
Time:                        23:17:28      BIC                      -63.836
Sample:                        0      HQIC                      -79.318
                                - 134
Covariance Type:                opg
=====

```

	coef	std err	z	P> z	[0.025	0.975]
intercept	7.8362	6.228	1.258	0.208	-4.370	20.042
drift	0.0018	0.001	1.512	0.131	-0.001	0.004
ar.L1	0.1448	0.477	0.303	0.762	-0.791	1.080
ar.L2	0.3471	0.519	0.669	0.504	-0.670	1.365
ar.L3	0.3695	0.332	1.114	0.265	-0.281	1.020
ma.L1	0.3589	0.469	0.765	0.444	-0.560	1.278
ma.L2	-0.1600	0.406	-0.395	0.693	-0.955	0.635
ma.L3	-0.3735	0.185	-2.020	0.043	-0.736	-0.011
sigma2	0.0261	0.004	6.699	0.000	0.018	0.034

```


=====
===
Ljung-Box (Q):                  41.37  Jarque-Bera (JB):
2.55
Prob(Q):                        0.41  Prob(JB):
0.28
Heteroskedasticity (H):        1.42  Skew:
-0.26
Prob(H) (two-sided):           0.24  Kurtosis:
2.56
=====
===

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-
step).
"""
```

```
[34]: forecasts = pd.Series(fit.predict(100), index=pd.period_range('2014',
    ↪periods=100, freq='Y'))
plt.figure()
df.plot()
forecasts.plot()
```

[34]: <matplotlib.axes.\_subplots.AxesSubplot at 0x12a1ae910>


<Figure size 432x288 with 0 Axes>



```
[35]: forecasts = pd.DataFrame(fit.predict(100, return_conf_int=True)[1],
    index=pd.period_range('2014', periods=100, freq='Y'))
plt.figure()
df.plot()
forecasts[0].plot()
forecasts[1].plot()
```

[35]: <matplotlib.axes.\_subplots.AxesSubplot at 0x12a3454d0>

<Figure size 432x288 with 0 Axes>



Note that forecast confidence intervals for the trend stationary model don't get bigger (in contrast with the difference stationary model).

[ ]: